

An Introduction to Deep Reinforcement Learning

JORGE C. CHAMBY DIAZ

JCHAMBYD@GMAIL.COM

JANUARY 08, 2020



Motivation

Can we create Artificial Intelligence?

Motivation

Can we create Artificial Intelligence?



Motivation


Can we create Artificial Intelligence?




We can create programs that LEARN!!

Deep Reinforcement Learning (Deep RL)

Deep Learning



Deep RL



- **What is it?** Framework for learning to solve sequential decision making problems.
- **How?** Trial and error in a world that provides occasional rewards
- **Deep?** Deep RL = RL + Neural Networks

Classes of Learning Problems

Supervised Learning

Data: (x, y)

x is data, y is label

Goal: Learn function to map
 $x \rightarrow y$

Apple example:



This thing is an apple.

Classes of Learning Problems

Supervised Learning

Data: (x, y)

x is data, y is label

Goal: Learn function to map
 $x \rightarrow y$

Apple example:



This thing is an apple.


Unsupervised Learning

Data: x

x is data, no labels!

Goal: Learn underlying
structure

Apple example:



This thing is like
the other thing.

Classes of Learning Problems

Supervised Learning

Data: (x, y)
 x is data, y is label

Goal: Learn function to map
 $x \rightarrow y$

Apple example:




This thing is an apple.

Unsupervised Learning

Data: x
 x is data, no labels!

Goal: Learn underlying
structure

Apple example:



This thing is like
the other thing.

Reinforcement Learning

Data: state-action pairs

Goal: Maximize future rewards
over many time steps

Apple example:



Eat this thing because it
will keep you alive.

Classes of Learning Problems

RL: our focus today!

Reinforcement Learning

Data: state-action pairs

Goal: Maximize future rewards over many time steps

Apple example:



Eat this thing because it will keep you alive.

Reinforcement Learning (RL): Key Concepts



Agent: take actions.

Reinforcement Learning (RL): Key Concepts



AGENT



ENVIRONMENT

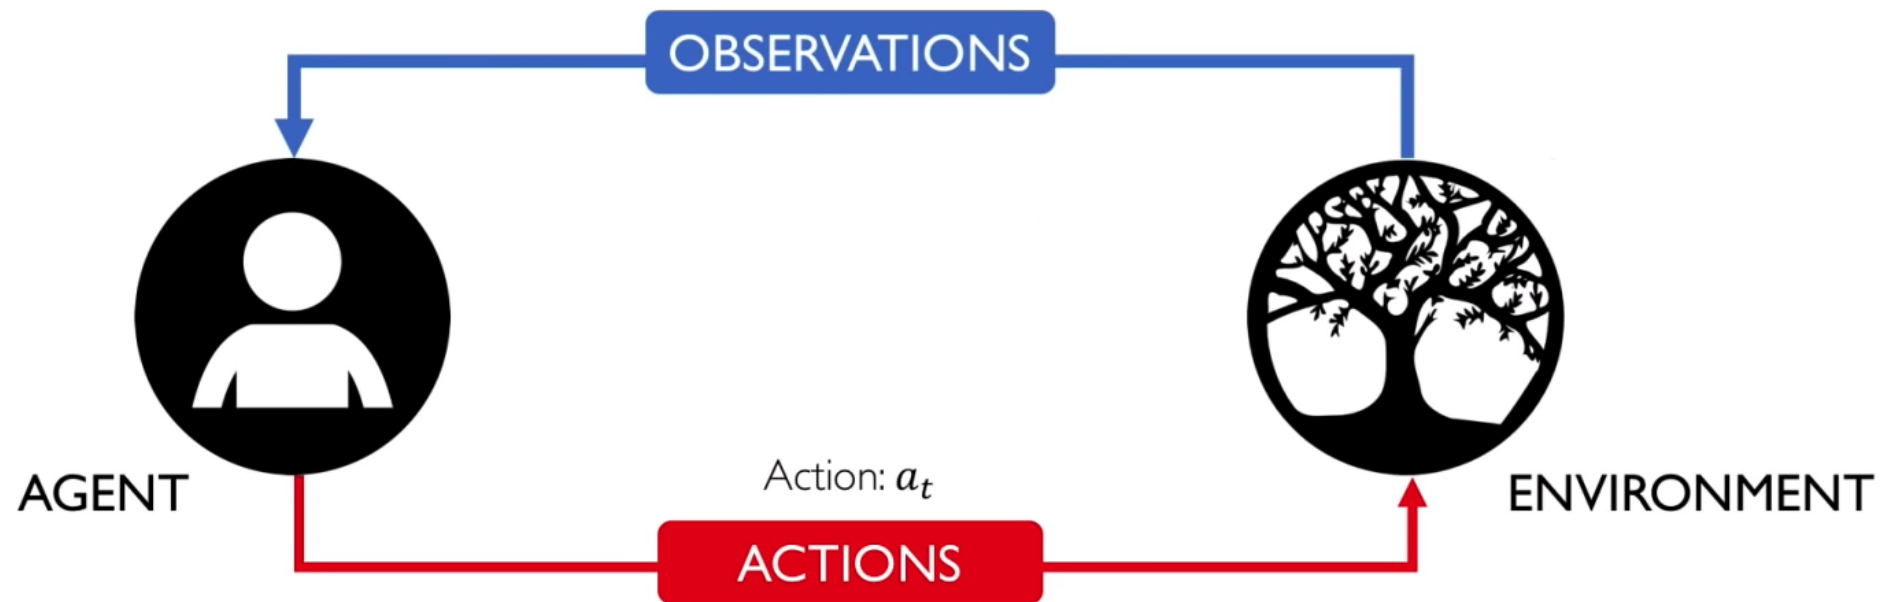
Environment: the world in which the agent exist and operates.

Reinforcement Learning (RL): Key Concepts



Action: a move the agent can make in the environment.

Reinforcement Learning (RL): Key Concepts



Observations: of the environment after taking actions.

Reinforcement Learning (RL): Key Concepts




State: an situation which the agent perceives.

Reinforcement Learning (RL): Key Concepts

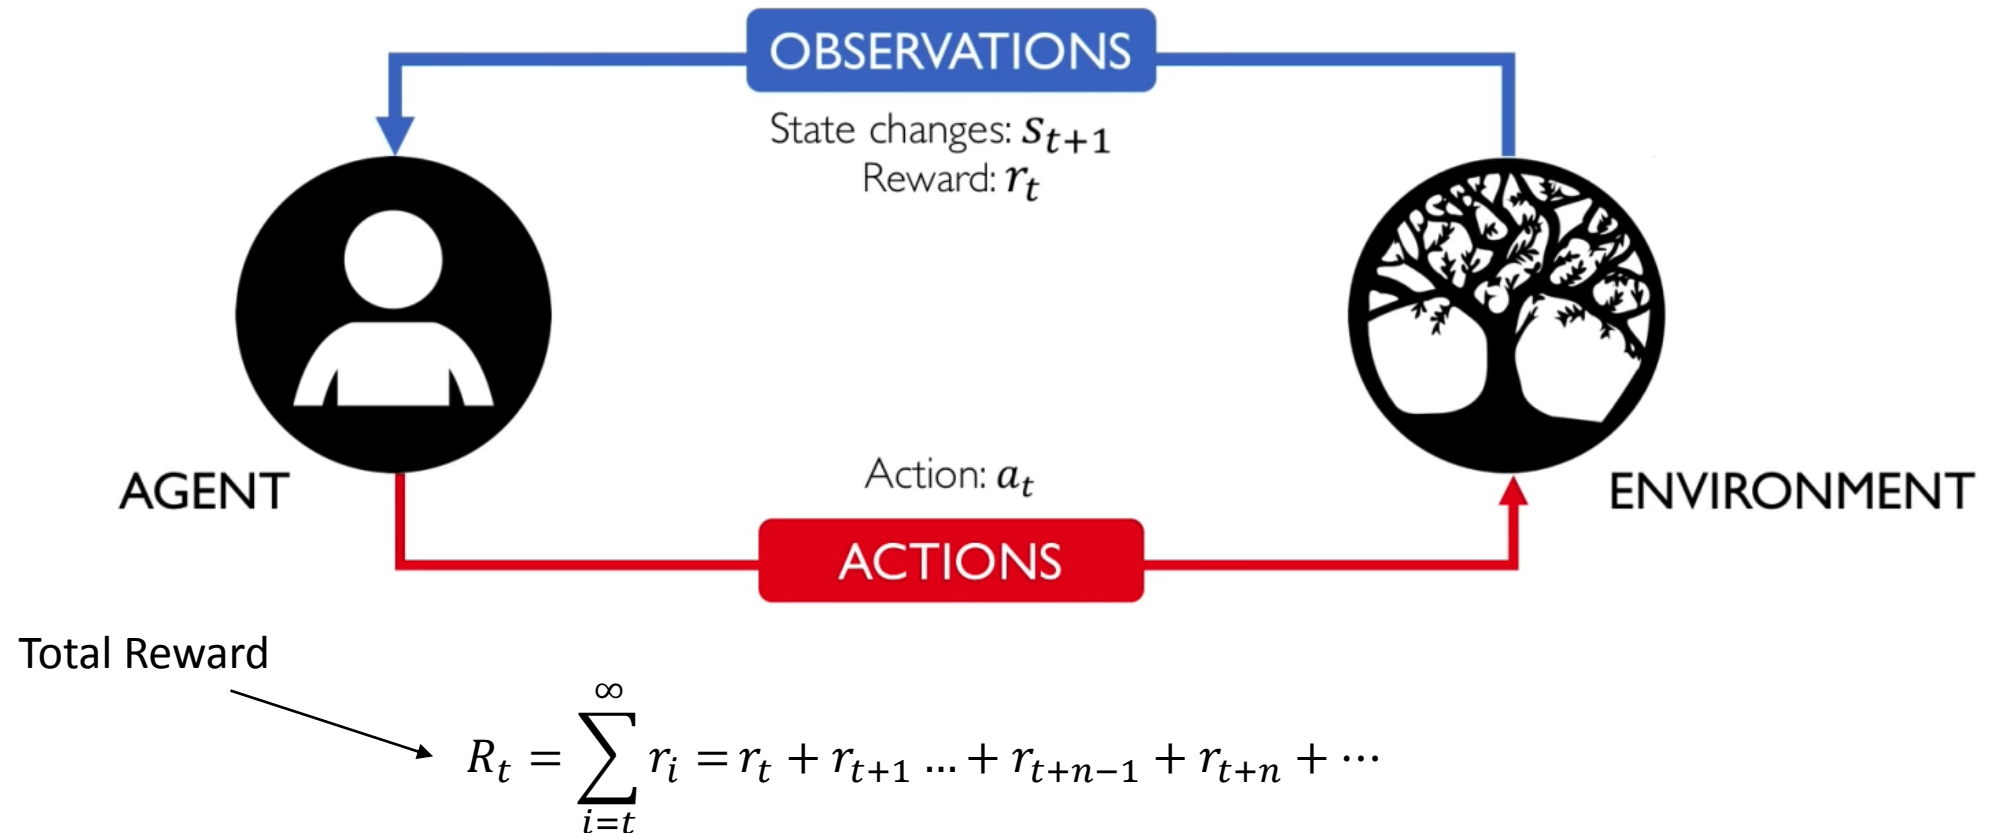


Reward: feedback that measure the success or failure of the agent's action.

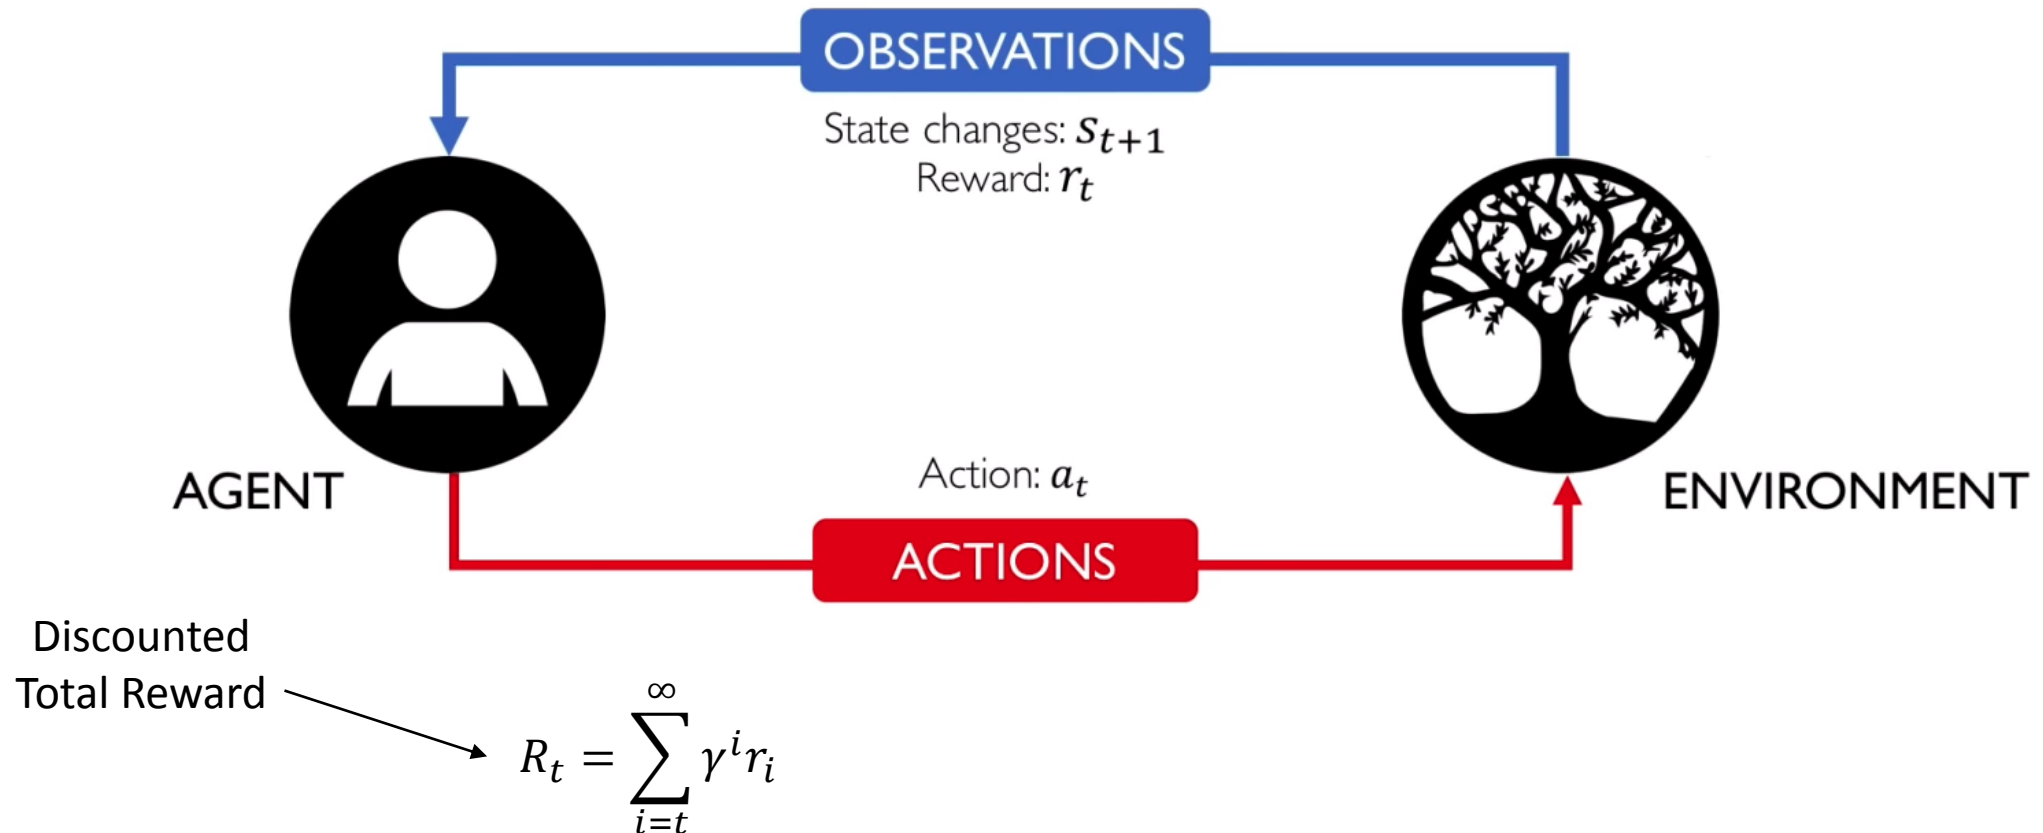
Reinforcement Learning (RL): Key Concepts



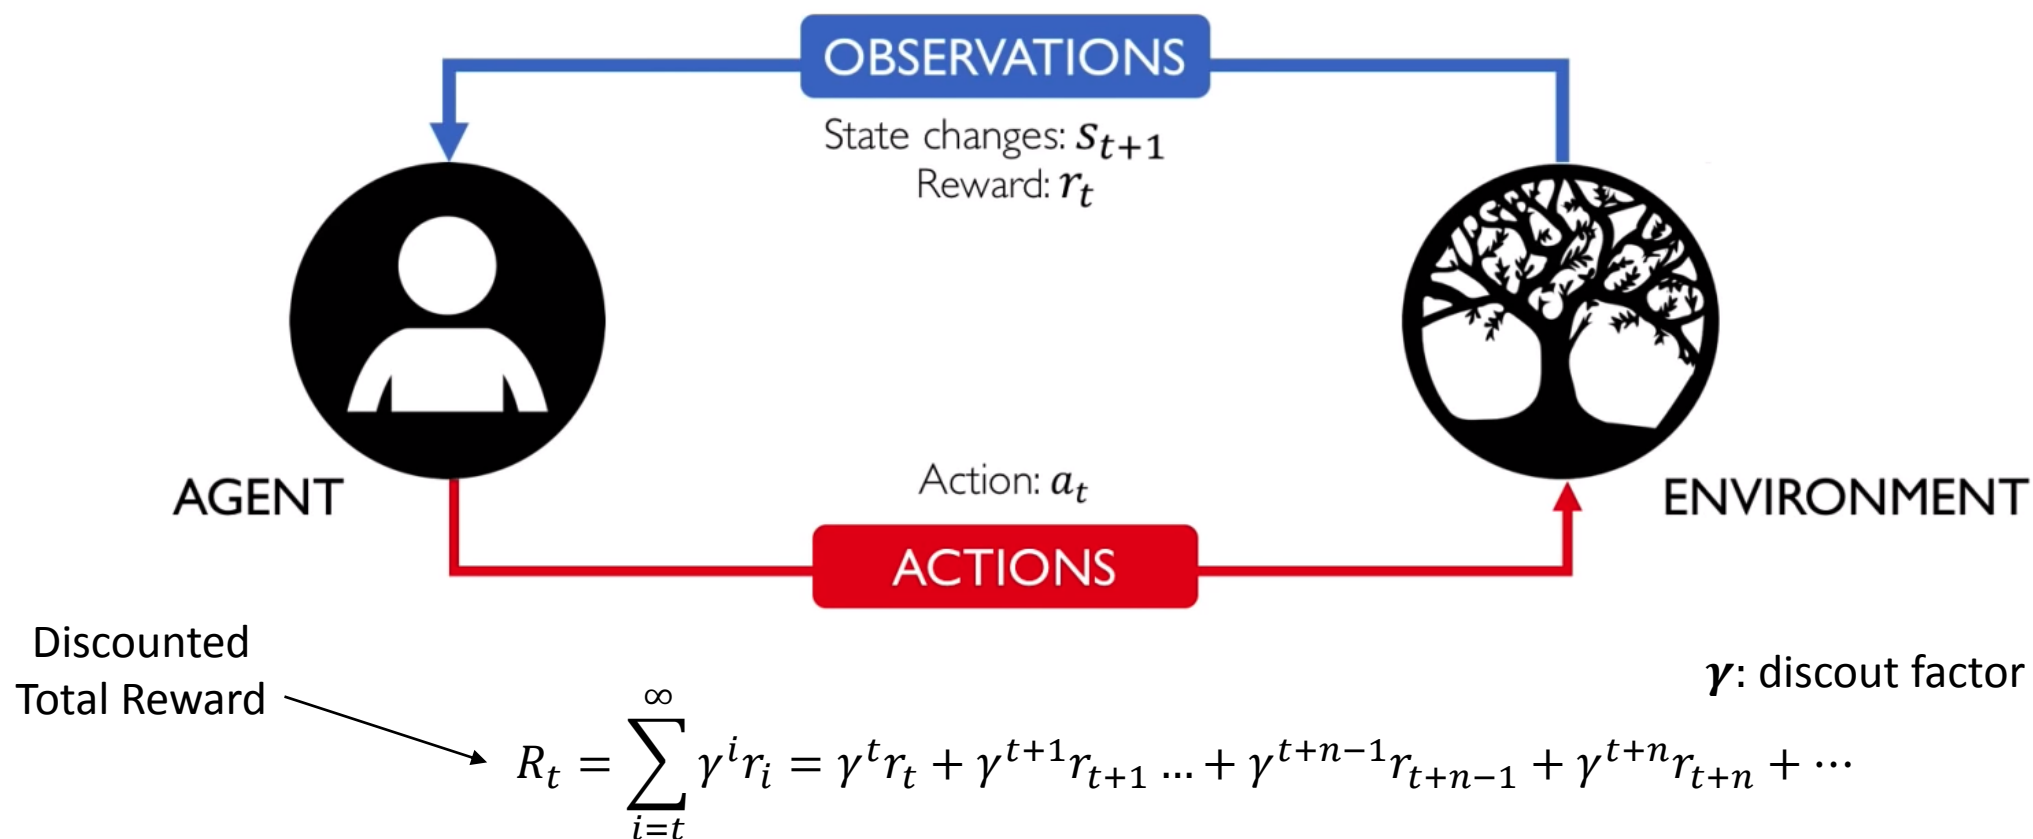
Reinforcement Learning (RL): Key Concepts




Reinforcement Learning (RL): Key Concepts



Reinforcement Learning (RL): Key Concepts



Example: Cart-Pole Problem




Objective: Balance a pole on top of a movable cart

State: angle, angular speed, position, horizontal velocity

Action: horizontal force applied on the cart

Reward: 1 at each time step if the pole is upright

Example: Robot Locomotion



Objective: Make the robot move forward

State: Angle and position of the joints

Action: Torques applied on joints

Reward: 1 at each time step upright + forward movement

Example: Atari games




Objective: Complete the game with the highest score

State: Raw pixel inputs of the game state

Action: Game controls e.g. Left, Right, Up, Down

Reward: Score increase/decrease at each time step

Example: Go



Objective: Win the game!

State: Position of all pieces

Action: Where to put the next piece down

Reward: 1 if win at the end of the game, 0 otherwise

How can we mathematically formalize the RL problem?



Markov Decision Problem

- Mathematical formulation of the RL problem
- **Markov property**: Current state completely characterises the state of the world

Defined by: $(\mathcal{S}, \mathcal{A}, \mathcal{R}, \mathbb{P}, \gamma)$

\mathcal{S} : set of possible states

\mathcal{A} : set of possible actions

\mathcal{R} : distribution of reward given (state, action) pair

\mathbb{P} : transition probability i.e. distribution over next state given (state, action) pair





γ : discount factor

Markov Decision Problem


- At time step $t=0$, environment samples initial state $s_0 \sim p(s_0)$
- Then, for $t=0$ until done:
 - Agent selects action a_t
 - Environment samples reward $r_t \sim R(\cdot | s_t, a_t)$
 - Environment samples next state $s_{t+1} \sim P(\cdot | s_t, a_t)$
 - Agent receives reward r_t and next state s_{t+1}
- A policy π is a function from S to A that specifies what action to take in each state
- **Objective:** find policy π^* that maximizes cumulative discounted reward: $\sum_{t \geq 0} \gamma^t r_t$

A simple MDP: Grid World

actions = {

1. right 
 2. left 
 3. up 
 4. down 
- }


states




Set a negative “reward”
for each transition
(e.g. $r = -1$)

Objective: reach one of terminal states (greyed out) in
least number of actions

A simple MDP: Grid World



Random Policy



Optimal Policy

Defining the Q-function

$$R_t = r_t + \gamma r_{t+1} + \gamma^2 r_{t+2} + \dots$$

Total reward, R_t , is the discounted sum of all rewards obtained from time t

$$Q(s, a) = \mathbb{E}[R_t]$$

The Q-function captures the **expected total feature reward** an agent in state, s , can receive by executing a certain **action, a**

How to take actions given a Q-function?

$$Q(\underset{\substack{\uparrow \\ (state)}}{s}, \underset{\substack{\uparrow \\ (action)}}{a}) = \mathbb{E}[R_t]$$

Ultimately, the agent needs a **policy** $\pi(s)$, to infer the **best action to take** at its state, s

How to take actions given a Q-function?

$$Q(\underset{\substack{\uparrow \\ (state)}}{s}, \underset{\substack{\uparrow \\ (action)}}{a}) = \mathbb{E}[R_t]$$

Ultimately, the agent needs a **policy** $\pi(s)$, to infer the **best action to take** at its state, s

Strategy: the policy should choose an action that maximizes futures reward

$$\pi^*(s) =$$

How to take actions given a Q-function?

$$Q(\underset{\substack{\uparrow \\ (state)}}{s}, \underset{\substack{\uparrow \\ (action)}}{a}) = \mathbb{E}[R_t]$$

Ultimately, the agent needs a **policy** $\pi(s)$, to infer the **best action to take** at its state, s

Strategy: the policy should choose an action that maximizes futures reward

$$\pi^*(s) = \underset{a}{\operatorname{argmax}} Q(s, a)$$

Deep Reinforcement Learning Algorithms

Value Learning

Find $Q(s, a)$

$$a = \underset{a}{\operatorname{argmax}} Q(s, a)$$

Policy Learning

Find $\pi(s)$

Sample $a \sim \pi(s)$

Digging deeper into the Q-function

Example: Atari Breakout




Digging deeper into the Q-function

Example: Atari Breakout




Digging deeper into the Q-function

Example: Atari Breakout





Digging deeper into the Q-function

Example: Atari Breakout



It can be very difficult for humans to accurately estimate Q-values




Which (s, a) pair has a higher Q-value?





Deep Q Networks (DQN): Training

How can we use deep neural networks to model Q-functions?



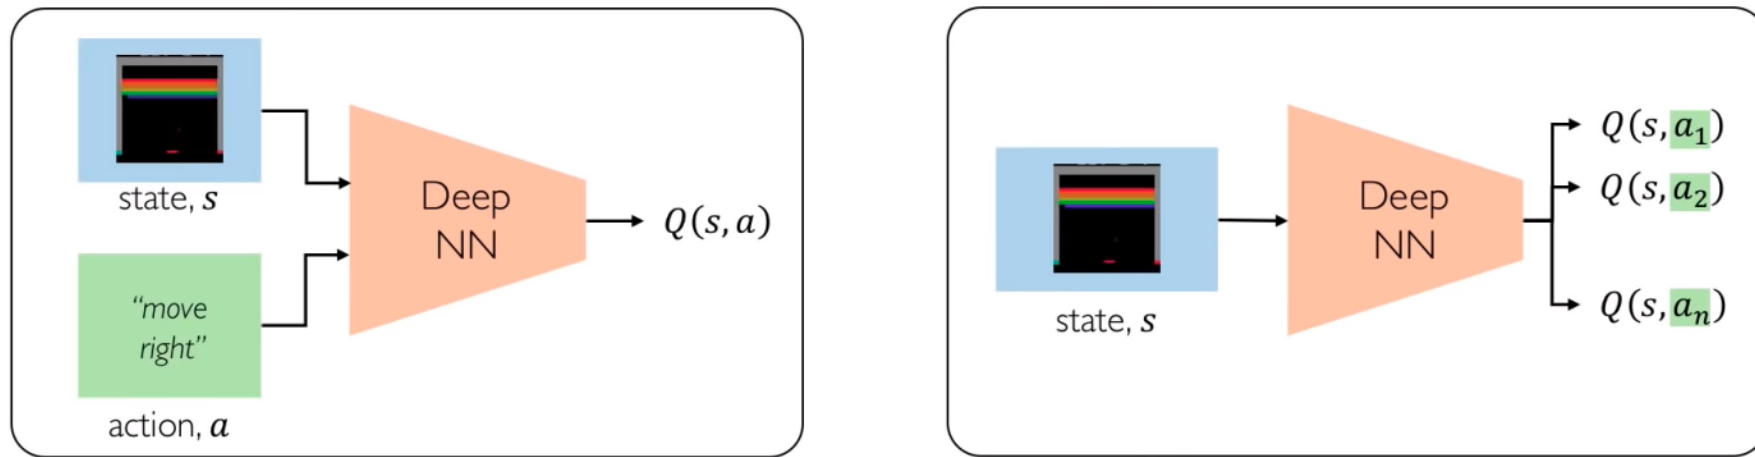
Deep Q Networks (DQN): Training

How can we use deep neural networks to model Q-functions?



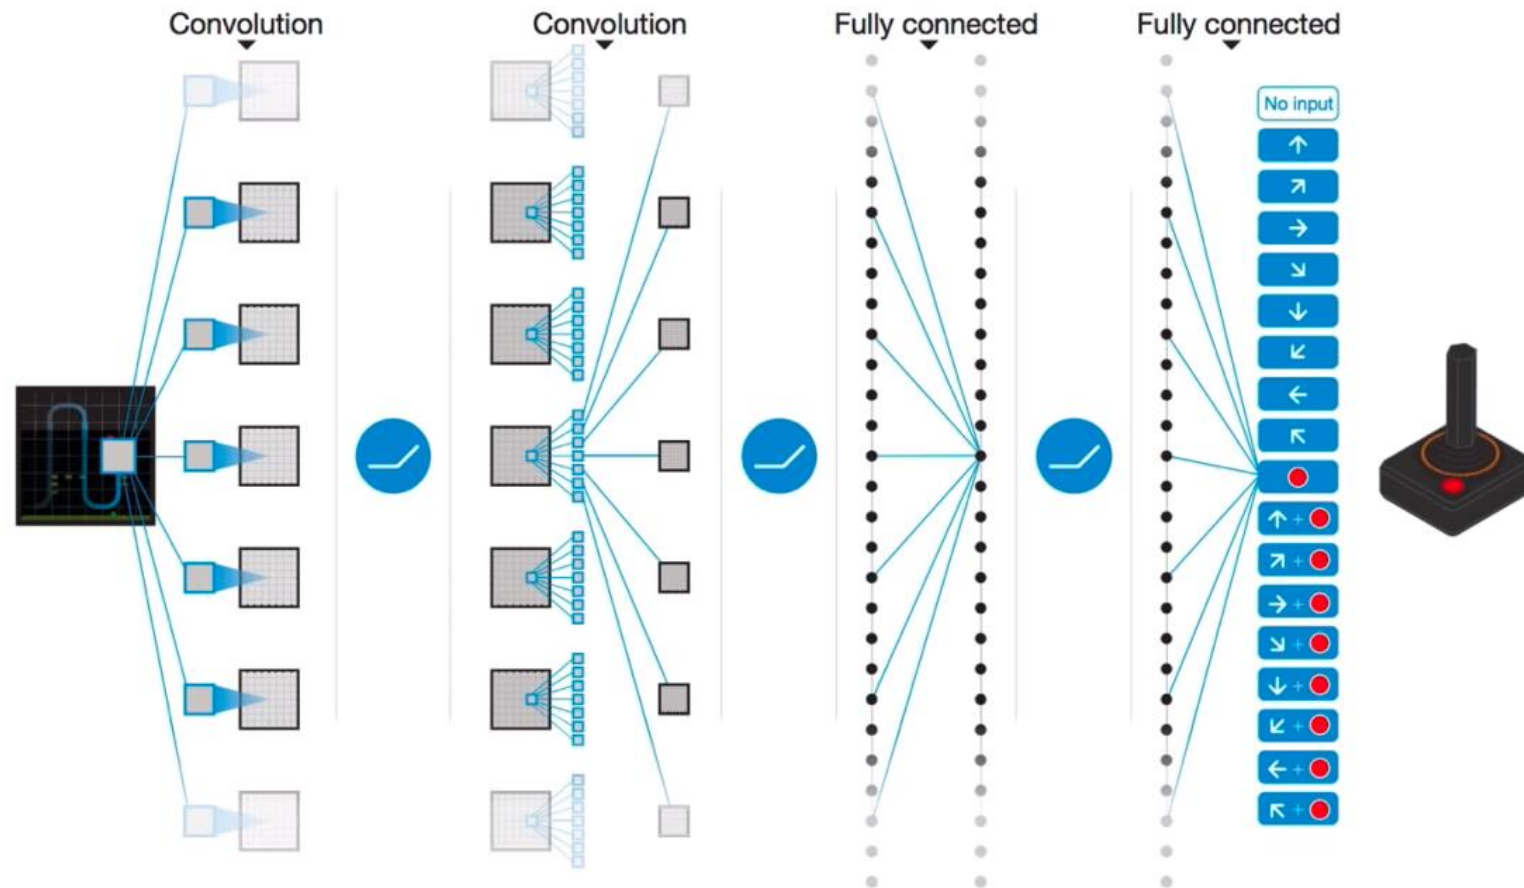
Deep Q Networks (DQN): Training

How can we use deep neural networks to model Q-functions?



$$\mathcal{L} = \mathbb{E} \left[\left\| \overbrace{\left(r + \gamma \max_{a'} Q(s', a') \right)}^{\text{target}} - \overbrace{Q(s, a)}^{\text{predicted}} \right\|^2 \right]$$

DQN Atari results



Downsides of Q-learning

Complexity:

- Can model scenarios where the action space is discrete and small
- Cannot handle continuous action spaces

Flexibility:

- Cannot learn stochastic policies since policy is deterministically computed from the Q function

Downsides of Q-learning

Complexity:

- Can model scenarios where the action space is discrete and small
- Cannot handle continuous action spaces

IMPORTANT:

Imagine you want to predict steering wheel angle of a car!

Flexibility:

- Cannot learn stochastic policies since policy is deterministically computed from the Q function

Downsides of Q-learning

Complexity:

- Can model scenarios where the action space is discrete and small
- Cannot handle continuous action spaces

IMPORTANT:

Imagine you want to predict steering wheel angle of a car!


Flexibility:

- Cannot learn stochastic policies since policy is deterministically computed from the Q function

**To overcome, consider a new class of RL training algorithms:
Policy gradient methods**

Policy Gradient (PG): Key Idea


DQN (before): Approximating Q and inferring the optimal policy,



Policy Gradient (PG): Key Idea

DQN (before): Approximating Q and inferring the optimal policy,


Policy Gradient: Directly optimize the policy!




Policy Gradient (PG): Key Idea

DQN (before): Approximating Q and inferring the optimal policy,

Policy Gradient: Directly optimize the policy!




Policy Gradient (PG): Training



1. Run a policy for a while
2. Increase probability of actions that lead to high rewards
3. Decrease probability of actions that lead to low/no rewards

```
function REINFORCE
  Initialize  $\theta$ 
  for  $episode \sim \pi_\theta$ 
     $\{s_i, a_i, r_i\}_{i=1}^{T-1} \leftarrow episode$ 
    for  $t = 1$  to  $T-1$ 
       $\nabla \leftarrow \nabla_\theta \log \pi_\theta(a_t | s_t) R_t$ 
       $\theta \leftarrow \theta + \alpha \nabla$ 
  return  $\theta$ 
```


Policy Gradient (PG): Training



1. Run a policy for a while
2. Increase probability of actions that lead to high rewards
3. Decrease probability of actions that lead to low/no rewards

```
function REINFORCE
  Initialize  $\theta$ 
  for  $episode \sim \pi_{\theta}$ 
     $\{s_i, a_i, r_i\}_{i=1}^{T-1} \leftarrow episode$ 
    for  $t = 1$  to  $T-1$ 
       $\nabla \leftarrow \nabla_{\theta} \log \pi_{\theta}(a_t | s_t) R_t$ 
       $\theta \leftarrow \theta + \alpha \nabla$ 
  return  $\theta$ 
```


log-likelihood of action

$$\nabla_{\theta} \log \pi_{\theta}(a_t | s_t) R_t$$

reward

The Game of Go

Aim: Get more board territory than your opponent.



Board Size $n \times n$	Positions 3^{n^2}	% Legal	Legal Positions
1×1	3	33.33%	1
2×2	81	70.37%	57
3×3	19,683	64.40%	12,675
4×4	43,046,721	56.49%	24,318,165
5×5	847,288,609,443	48.90%	414,295,148,741
9×9	$4.434264882 \times 10^{38}$	23.44%	$1.03919148791 \times 10^{38}$
13×13	$4.300233593 \times 10^{80}$	8.66%	$3.72497923077 \times 10^{79}$
19×19	$1.740896506 \times 10^{172}$	1.20%	$2.08168199382 \times 10^{170}$

Greater number of legal board positions than atoms in the universe.


Source: Wikipedia.

AlphaGo Beats Top Human Player at Go (2016)



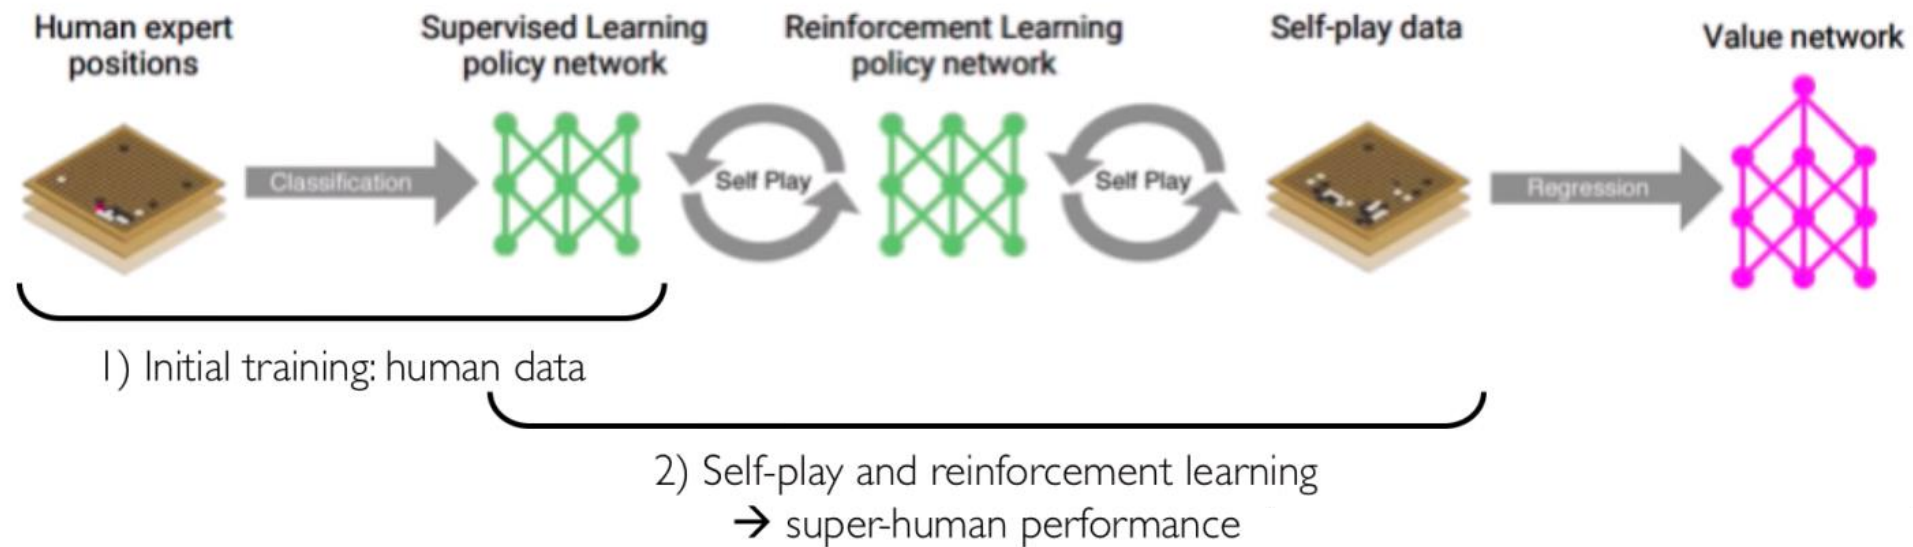
Silver et al., *Nature* 2016.

AlphaGo Beats Top Human Player at Go (2016)



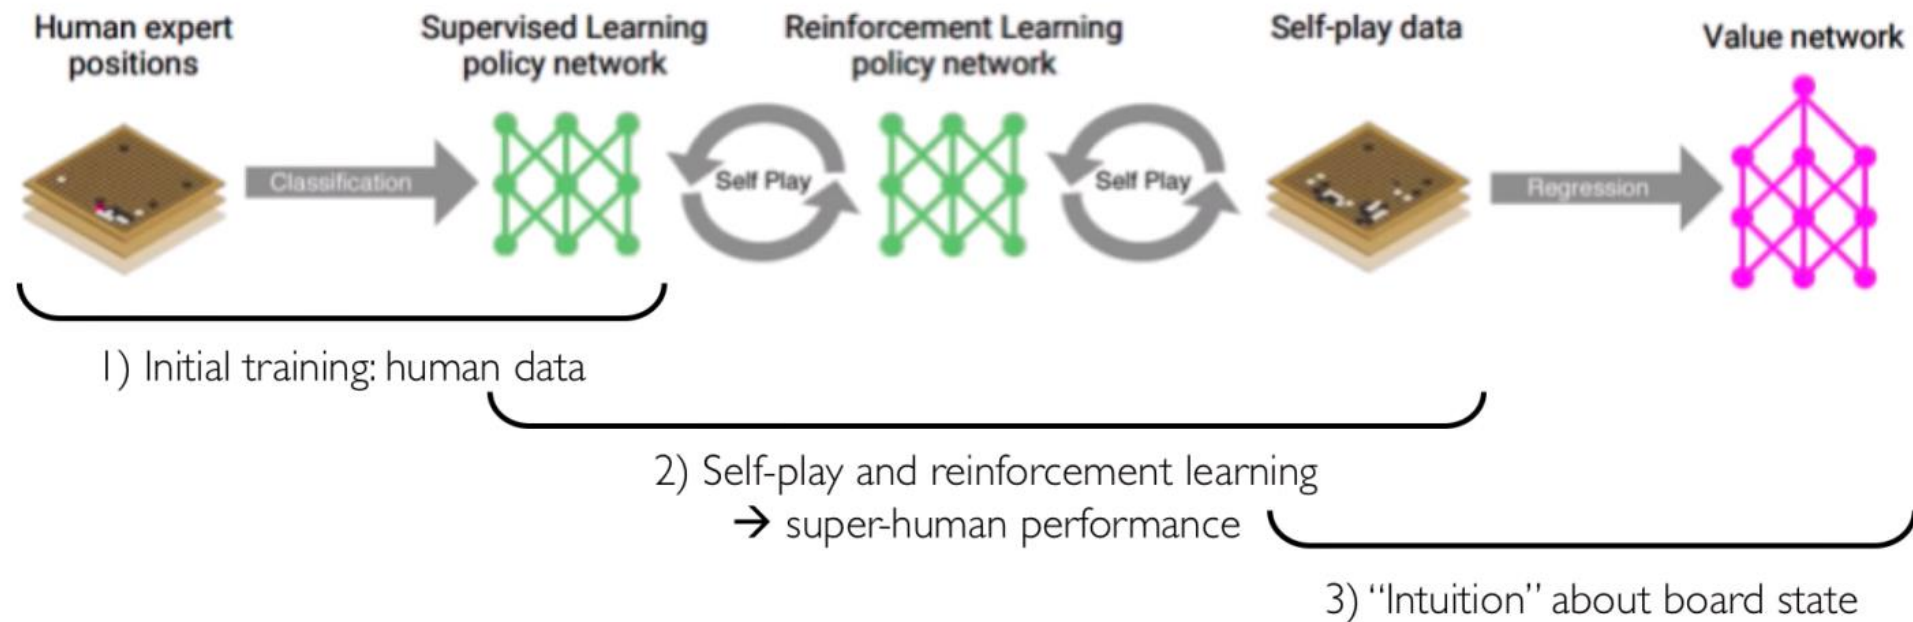
Silver et al., *Nature* 2016.

AlphaGo Beats Top Human Player at Go (2016)



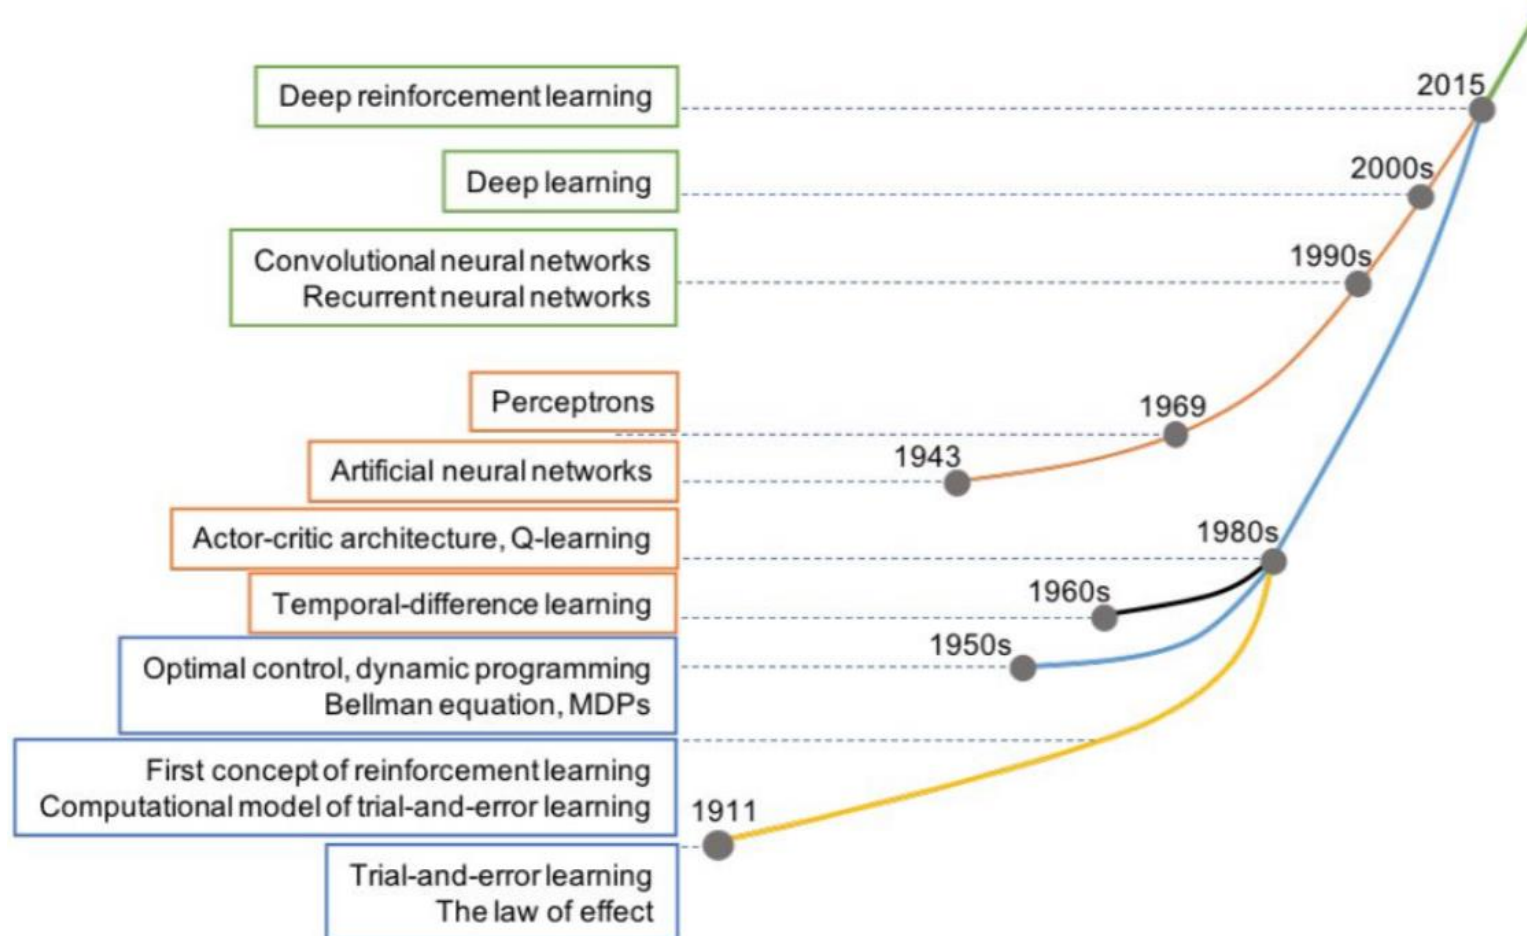
Silver et al., *Nature* 2016.

AlphaGo Beats Top Human Player at Go (2016)



Silver et al., *Nature* 2016.

RL Milestones



Questions?